Difference Between RDD, DataFrame, and Dataset in Apache Spark

Apache Spark provides **three main data abstractions** for processing data:

1. **RDD (Resilient Distributed Dataset)** – Low-level API
2. **DataFrame** – High-level API
3. **Dataset** – High-level, strongly typed API (Scala/Java only)

Each abstraction offers a different balance between **performance, ease of use, and type safety**.

Apache Spark Data Abstractions

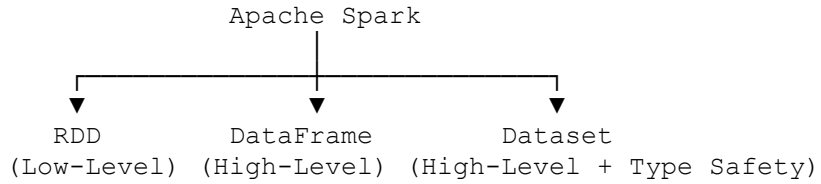

---

1. RDD (Resilient Distributed Dataset)

What is an RDD?

An **RDD** is Spark's original distributed data structure. It is an **immutable**, **fault-tolerant**, distributed collection of objects.

RDDs provide low-level control over distributed computations.

Characteristics

- Low-level API
- Stores any type of object
- Immutable
- Distributed across partitions
- Fault tolerant through lineage
- No schema
- No Catalyst Optimizer
- No Tungsten optimization

Example

```
rdd = spark.sparkContext.parallelize([1, 2, 3, 4])

result = rdd.map(lambda x: x * 2)

print(result.collect())
```

Advantages

- Fine-grained control
- Suitable for complex custom transformations
- Works with unstructured data

Disadvantages

- Slower than DataFrames and Datasets
- No automatic query optimization
- No schema
- More code to write

---

2. DataFrame

What is a DataFrame?

A **DataFrame** is a distributed collection of data organized into **named columns**, similar to a table in a relational database or a Pandas DataFrame.

Internally, a DataFrame is a **Dataset of Rows**.

Characteristics

- High-level API
- Has a schema
- Uses Catalyst Optimizer
- Uses Tungsten execution engine
- Much faster than RDD
- Supports SQL queries

Example

```
df = spark.read.csv("employees.csv", header=True)

df.filter(df.salary > 5000).show()
```

Advantages

- Very fast
- Easy to use
- SQL support
- Automatic optimization
- Less memory usage
- Preferred for ETL and analytics

Disadvantages

- Rows are not strongly typed
- Less low-level control than RDD

---

3. Dataset

What is a Dataset?

A **Dataset** combines the advantages of **RDDs** and **DataFrames**.

It provides:

- Strong compile-time type safety
- Catalyst optimization
- Tungsten execution
- Object-oriented programming support

Important: The Dataset API is **available only in Scala and Java**. **PySpark does not support typed Datasets.**

Scala Example

```
case class Employee(name: String, salary: Double)

val ds = Seq(
  Employee("Ali", 5000),
  Employee("Sara", 7000)
).toDS()
```

Advantages

- Strongly typed
- Compile-time error checking

- High performance
- Catalyst optimization
- Object-oriented programming

Disadvantages

- Not available in Python
- Slightly more complex than DataFrames

Performance Comparison

```
Performance

RDD
Dataset
DataFrame
```

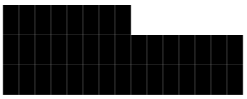

```
Fastest → DataFrame ≈ Dataset
Slowest → RDD
```

DataFrames and Datasets are significantly faster because Spark optimizes their execution plans using **Catalyst** and **Tungsten**.

Feature Comparison

Feature	RDD	DataFrame	Dataset
API Level	Low	High	High
Schema	❌ No	✅ Yes	✅ Yes
Type Safety	❌ No	❌ No (Rows)	✅ Yes
Catalyst Optimizer	❌ No	✅ Yes	✅ Yes
Tungsten Engine	❌ No	✅ Yes	✅ Yes
SQL Support	❌ No	✅ Yes	✅ Yes
Compile-time Type Checking	❌ No	❌ No	✅ Yes
Performance	Good	Excellent	Excellent
Languages	Scala, Java, Python	Scala, Java, Python, R	Scala, Java

---

When Should You Use Each?

Use RDD When

- You need low-level transformations.
- You are working with unstructured data.
- You need custom distributed algorithms.
- You require direct control over partitions and execution.

Use DataFrame When

- Building ETL pipelines.
- Working with structured or semi-structured data.
- Running SQL queries.
- Performing joins, aggregations, and analytics.
- Using PySpark.

This is the most common choice for Data Engineering.

Use Dataset When

- Using Scala or Java.
- You want compile-time type safety.
- You need object-oriented programming with Spark.
- You want DataFrame performance with typed objects.

Relationship Between Them

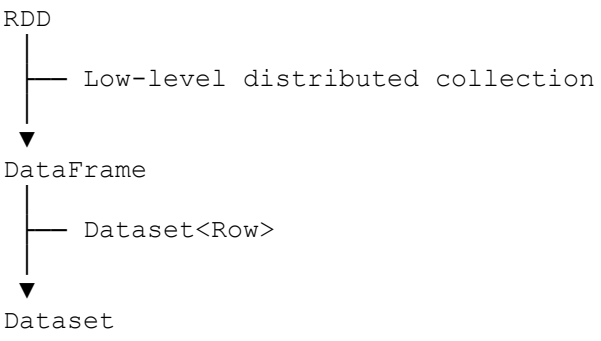

- **DataFrame** is essentially a **Dataset of Row objects**.
- **Dataset** extends the DataFrame concept by adding strong typing.
- **RDD** is the foundational low-level abstraction.

Interview Questions

Q1. Which is the fastest: RDD, DataFrame, or Dataset?

DataFrame and Dataset are generally faster than RDD because Spark optimizes them using the **Catalyst Optimizer** and **Tungsten execution engine**.

Q2. Why is DataFrame faster than RDD?

Because DataFrames have:

- Schema information
- Catalyst Optimizer
- Tungsten execution engine
- Better memory management
- Optimized query execution

RDDs do not have these optimizations.

Q3. Can you use SQL with an RDD?

No.

SQL queries work with **DataFrames** and **Datasets**, not directly with RDDs.

Q4. Does PySpark support Dataset?

No.

PySpark supports:

- RDD
- DataFrame

Typed **Dataset** is available only in **Scala and Java**.

Q5. Which abstraction is most commonly used in Data Engineering?

DataFrame is the standard choice because it offers high performance, SQL support, automatic optimization, and ease of use.

Interview Summary

- **RDD** is Spark's original low-level distributed data abstraction. It is immutable, fault-tolerant, and provides fine-grained control, but it does not support schemas or automatic query optimization.
- **DataFrame** is a high-level distributed table with named columns. It supports SQL, uses the **Catalyst Optimizer** and **Tungsten** execution engine, and is the **most commonly used API in modern Spark applications**, especially in PySpark.
- **Dataset** combines the performance of DataFrames with **strong compile-time type safety**, but it is available only in **Scala and Java**.
- In modern Spark development:
    - **Use DataFrames** for almost all ETL, analytics, and SQL workloads.
    - **Use RDDs** only when low-level distributed control is required.
    - **Use Datasets** when working in Scala or Java and type safety is important.